

ROHIT TANEJA

AI Compute Infrastructure | Ex-NVIDIA, Apple | 15+ Years Scaling GPU & ML Platforms

rohit23taneja@gmail.com | +1-919-675-9274 | San Francisco Bay Area | US Permanent Resident

[LinkedIn](#) | [GitHub](#) | [Google Scholar](#)

VISION

The great problems in AI won't be solved by technology alone. They'll be solved by people: the culture, the teams, the structure around the technology. Infrastructure is the foundation. People are the multiplier. Get both right, and magic happens. I want to build that: infrastructure and teams built for a mission, with people who treat the mission as the boss and never stop adapting.

PROFESSIONAL EXPERIENCE

Independent Research: Agentic AI Security, Autonomy & Governance

May 2026 – Present | San Francisco, CA

- Completed Harvard Business School Executive MBA (July 2026)
- Co-authored "**The Digital Apprentice: A Framework for Human-Directed Agentic AI Development**," accepted at the **ACM AI Leadership Summit** (2026); research on balancing human oversight with autonomy in agentic AI systems (arXiv:2606.04321)
- Building an inference-time control plane for agentic AI systems: a multi-policy gateway that routes and evaluates production traffic across competing serving policies, with runtime drift detection that automatically recovers quality when a single policy degrades

Vice President, AI/ML Software | Lemurian Labs

January 2026 – May 2026 | San Francisco, CA

Led design of AI/ML serving architecture for heterogeneous compute fleets.

- Designed the AI/ML serving architecture spec spanning compiler through accelerator integration for heterogeneous hardware fleets
- Designed the multi-tenant serving engine spec, defining workload isolation and scheduling controls for shared compute infrastructure

Lead/Staff Member of Enterprise Product Group | NVIDIA

September 2021 – December 2025 | Santa Clara, CA

Owned DGX Cloud and NVIDIA Inference Microservices (NIM), NVIDIA's managed GPU compute and inference platform.

- Founding team for DGX Cloud; took it from MVP to production, defining the isolated compute/storage substrate and data-integration model for **top hyperscalers and frontier labs**
- Scaled the platform to **1,000+ fine-tuning jobs**, standardizing the GPU-consumption API spec for managed-service mode and building metered-usage billing integration with Hugging Face
- Partnered directly with **data platforms managing exabytes of data** on Day-2 operations: data manageability, ingress/egress, and staging workflows for GPU fleets running at hyperscaler scale
- Built a **Mistral-NVIDIA joint model** on CoreWeave infrastructure, leading joint training and data-sampling recipe setup
- Onboarded **10+ model and data partners** onto NIM; contributed to TensorRT-LLM optimization work for Meta Llama 3 (405B)
- Hold a patent on workload characterization and heterogeneous-fit auto-routing across NVIDIA GPU fleets

Lead, ML Infrastructure | Apple

October 2019 – September 2021 | Cupertino, CA

Built the real-time inference infrastructure for Apple's global manufacturing operations.

- Deployed deep learning inference at manufacturing scale, achieving **\$50M+ in measurable CapEx savings** and yield improvements on the iPhone line
- Built a **real-time inference platform spanning on-prem and cloud (CSP) infrastructure**, scaling from a single line of business to all Apple hardware products
- Designed a centralized data pipeline (Spark streaming) for continuous inference and drift detection
- Led global infrastructure delivery for always-on deployments across manufacturing sites in **India and China**, covering operating and reliability at physical-facility scale, not just software

Server Platform Performance Engineer, Big Data | IBM

October 2014 – October 2016 | Austin, TX

- Built a **distributed data processing pipeline** on Apache Spark over a Hadoop data lake, optimizing shuffle/sort performance at scale
- Containerized the platform (Spark-on-Docker) for enterprise cloud deployment
- Awarded **3x Eminence & Excellence Awards** for systems-performance contributions to National Labs and Enterprise customers

TECHNICAL EXPERTISE

COMPUTE INFRASTRUCTURE

GPU fleet architecture, DGX Cloud, NIM, multi-tenant serving, hyperscaler integration patterns (AWS, GCP, Azure)

HARDWARE & MANUFACTURING

GPU-accelerated architecture design, AI server PCB quality (AOI), compiler/hardware optimization (LLVM), factory-scale deployment

DISTRIBUTED SYSTEMS

Apache Spark, Hadoop, in-memory processing, Docker, Kubernetes

RELIABILITY & OPS

Drift detection and observability for always-on inference, Day-2 operations at hyperscaler scale, global infrastructure delivery

EDUCATION

Harvard Business School, Executive MBA, 2024–2026 (Completed July 2026)
North Carolina State University, M.S., Computer Engineering, 2011–2013